

Evidence-Free Decision Points in Biomarker-Driven Metastatic Cancer Guidelines: A Pre-Registered Multi-Tumor Audit of ESMO, ASCO, and NCCN Decision Trees in HR+/HER2- Breast Cancer and EGFR/ALK Non-Small-Cell Lung Cancer

H.-T. Lin^{1,*}

¹[Affiliation to be supplied at submission].

*Correspondence: ppoi87@gmail.com.

May 16, 2026

Abstract

Purpose. Guidelines for biomarker-driven metastatic cancers compose treatment recommendations across non-overlapping pivotal trials. The structural concordance between these decision trees and the underlying trial evidence has not been quantified at multi-tumor scale. **Methods.** We applied a pre-registered graph-theoretic framework to two biomarker-driven mBC subtypes (HR+/HER2-) and NSCLC subtypes (EGFR-mutant, ALK-rearranged), using systematic ClinicalTrials.gov searches (NSCLC freeze 2026-05-16, 702 candidates filtered to 180 pivotal trials; mBC corpus 81 trials from v2.0.0). The unified ESMO+ASCO+NCCN decision tree comprised 50 nodes (25 mBC + 25 NSCLC). The evidence-free decision-point ratio (EFDPR) is the fraction of guideline nodes with no trial edge satisfying state-superset, biomarker-concordance, drug-class, and temporal-precedence criteria; the pre-registered primary test is the one-sided exact binomial test of $H_0 : \text{EFDPR} \leq 0.25$. **Results.** Under strict concordance the pooled EFDPR was 0.48 (24/50; Clopper-Pearson 95% CI 0.34–0.63); the pre-registered exact-binomial test *rejected* H_0 at $\alpha = 0.05$ ($P = 0.0004$). Tumor-stratified sensitivity: mBC-only 0.40 ($P = 0.07$, marginal); NSCLC-only 0.56 ($P = 0.0009$, reject); NSCLC EGFR-only 0.71 ($P = 0.0001$, strong rejection); NSCLC ALK-only 0.25 ($P = 0.63$, well-evidenced). The largest evidence gaps clustered at post-osimertinib EGFR-mutant decision nodes (HER3-ADC, TROP2-ADC, amivantamab+chemo positions) and at post-CDK4/6i mBC decision nodes (everolimus salvage, no-actionable-mutation). **Conclusion.** At the unified ESMO+ASCO+NCCN scale across two biomarker-driven solid tumors, nearly half of guideline decision nodes lack directly-supporting pivotal trial evidence chains. The gap is largest in NSCLC EGFR-mutant post-osimertinib settings (71% evidence-free). Guideline committees should consider routine EFDPR audits to identify which decision recommendations rest on direct trial evidence and which on cross-trial extrapolation.

Key Objective. To quantify, at pooled mBC+NSCLC scale, how many ESMO/ASCO/NCCN guideline decision nodes lack directly-supporting pivotal trial evidence chains.

Knowledge Generated. On 50 unified decision nodes against a 210-edge trial-DAG drawn from 261 systematically-searched pivotal trials, the strict EFDPR is 0.48 (Clopper-Pearson 95% CI 0.34–0.63). The pre-registered one-sided exact-binomial test rejects $H_0 : \text{EFDPR} \leq 0.25$ at $P = 0.0004$. NSCLC EGFR-mutant post-osimertinib decisions are the most evidence-sparse (71% evidence-free).

Relevance. The EFDPR overlay locates the specific decision nodes — primarily NSCLC EGFR-mutant post-osimertinib and mBC post-CDK4/6i salvage — where current evidence chains rest on extrapolation rather than direct pivotal-trial support, providing both an audit tool for guideline committees and a target-population specification for trial sponsors.

Keywords: metastatic breast cancer; non-small-cell lung cancer; EGFR; ALK; treatment sequencing; clinical guidelines; computational meta-research; graph-theoretic framework.

1 Introduction

Modern systemic therapy for biomarker-driven metastatic cancers has expanded dramatically over the past decade: in HR+/HER2-negative metastatic breast cancer (mBC), CDK4/6 inhibitors, selective estrogen receptor degraders, PI3K/AKT-pathway inhibitors, HER2-low-targeted antibody-drug conjugates, and TROP2-directed ADCs have all entered routine practice;¹⁻³ in EGFR-mutant and ALK-rearranged non-small-cell lung cancer (NSCLC), third-generation EGFR-tyrosine kinase inhibitors, bispecific antibodies (amivantamab + lazertinib), HER3-directed ADCs, third-generation ALK inhibitors, and TROP2-ADCs have followed similar trajectories.^{4,5} Each agent achieved approval on the basis of a pivotal trial that fixed a single point in the multi-line trajectory, typically without head-to-head testing against simultaneously-approved alternatives.

Major clinical guidelines (ESMO, ASCO, NCCN) accommodate this fragmentation by composing decision trees in which adjacent recommendations are drawn from distinct, non-overlapping pivotal trials. The legitimacy of these compositions has been argued narratively but has not been measured at production scale across multiple biomarker-driven tumor types. Our group’s earlier single-tumor mBC analysis (v2.0.0; preserved tag) found a marginal non-rejection ($P = 0.07$) at $n = 25$ guideline nodes — consistent with substantial composition-across-non-overlapping-trials but underpowered for confirmatory inference. The present manuscript extends to NSCLC (EGFR + ALK), pools the two-tumor decision-tree denominator to $n = 50$, and applies the pre-registered exact-binomial test of $H_0 : \text{EFDPR} \leq 0.25$ for the first time at adequately-powered scale.

2 Methods

2.1 Data sources and pre-registration

The pre-registration (docs/prereg-v3.md, commit 4b5bf1a) was committed before any NSCLC outcome-touching analysis. The primary outcome (pooled strict-tolerance EFDPR), primary test (one-sided exact binomial of $H_0 : p \leq 0.25$ at $\alpha = 0.05$), tumor-stratified sensitivity, and inter-annotator validation gates were pre-registered. mBC analysis follows v2.0.0 (frozen tag); NSCLC analysis is new in v3.

2.2 Systematic trial search and inclusion

NSCLC: ClinicalTrials.gov v2 API (freeze 2026-05-16; 702 candidates) was filtered by eight pre-specified criteria (phase 2/3, interventional, start year 2013–2026, NSCLC + metastatic, EGFR or ALK biomarker, eligibility retrievable, enrolment ≥ 100 , primary completion by 2026) to 176 systematic candidates. Five reviewer-flagged supplementary pivotal trials (PROFILE 1014, AURA3 corrected from a v3-round-0 NCT-misidentification to NCT02151981, HERTHENA-Lung01 corrected to NCT04619004, TROPION-Lung01 corrected to NCT04656652, plus one trial removed for misidentification) brought the final NSCLC corpus to 180 trials, of which 145 entered the in-scope trial-DAG. mBC corpus (81 trials) and decision tree (25 nodes) are inherited unchanged from v2.0.0.

2.3 Dual-annotator LLM extraction

For NSCLC, eligibility text was structured into the schema (v1.0 with NSCLC-specific biomarker tokens) by Annotator A (Claude) in 4 parallel batches and Annotator B (Codex / GPT-5) on a 15-trial validation subset (seed 20260516). Pre-adjudication mean Cohen’s κ on 8 NSCLC key fields was 0.78; mean PABAK was 0.73. The kappa gate failed pre-adjudication on `post_alk_tki` (small validation-sample effect, $\kappa = 0.29$) and `egfr_t790m` ($\kappa = 0.61$); other fields passed.




Figure 1. Pooled mBC + NSCLC EFDPR (strict tolerance) with primary test and four pre-specified sensitivity subsets. Diamond: primary 50-node analysis ($P = 0.0004$); circles: sensitivity subsets. Asterisks mark rejections at $\alpha = 0.05$.

2.4 NSCLC guideline decision-tree encoding

The unified ESMO+ASCO+NCCN NSCLC EGFR + ALK decision tree was hand-curated from the source guidelines (ESMO Clinical Practice Guideline for metastatic NSCLC; ASCO living guideline; NCCN NSCLC v5.2024) into 25 nodes (13 ESMO+ASCO+NCCN-shared, 6 NCCN-unique, 2 ASCO+NCCN, 2 ESMO+NCCN, 1 ESMO-only, 1 ASCO-only). 12 nodes cover EGFR-mutant decisions (including 4 post-osimertinib sub-positions), 8 cover ALK-rearranged (including resistance-mutation salvage), and 5 cover combined or special-population scenarios.

2.5 Concordance algorithm and pooled inference

The four concordance criteria are unchanged from v2: state superset, drug-class equivalence, biomarker concordance (pre-registered tolerance grid: strict / ESCAT / liberal), and temporal precedence. The combined DAG has 210 edges (145 NSCLC in-scope + 65 mBC in-scope) and 47 canonical source nodes. The pooled EFDPR is computed on the 50-node pooled denominator with the pre-registered one-sided exact binomial test as primary; tumor-stratified estimates (mBC, NSCLC, NSCLC-EGFR, NSCLC-ALK) are reported as sensitivity.

3 Results

3.1 Pooled primary test rejects the null

Under strict concordance the pooled EFDPR was **0.480** (24/50 decision nodes evidence-free; Clopper-Pearson 95% CI 0.34–0.63). The pre-registered one-sided exact binomial test of $H_0 : \text{EFDPR} \leq 0.25$ **rejected the null** at $\alpha = 0.05$ with $P = 0.0004$ (Figure 1). ESCAT-aligned and liberal tolerance gave 0.34 ($P = 0.098$), indicating that some of the strict-tolerance gap reflects biomarker-operationalization stringency, but the strict result is the pre-registered primary inference.




Figure 2. Pilot → production → multi-tumor trajectory of EFDPR (left axis) and the pre-registered exact-binomial test P -value (right axis, $-\log_{10}$ scale). The pooled 50-node test rejects at $P = 0.0004$ ($-\log_{10} P = 3.4$).

3.2 Pilot-to-multi-tumor trajectory of the primary inference

The pre-registered exact-binomial test rejected the null on the v3 pooled ($n=50$) corpus despite failing to reject on v1.0.0 (mBC pilot, $n=16$, $P = 0.37$) and v2.0.0 (mBC production, $n=25$, $P = 0.07$). The trajectory is consistent with the power gain from doubling the guideline-node denominator and adding a tumor (NSCLC) whose post-osimertinib evidence base is sparser than mBC's post-CDK4/6i evidence base (Figure 2).

3.3 Tumor-stratified sensitivity

NSCLC-only ($n = 25$) gave EFDPR 0.56 ($P = 0.0009$, rejects); NSCLC EGFR-only ($n = 17$) gave 0.71 ($P = 0.0001$, strongly rejects); NSCLC ALK-only ($n = 8$) gave 0.25 ($P = 0.63$, fails to reject); mBC-only ($n = 25$) gave 0.40 ($P = 0.07$, marginal). The pooled rejection is driven by NSCLC EGFR-mutant decision nodes, particularly post-osimertinib subpositions; ALK-rearranged decisions are comparatively well-evidenced.

3.4 Where the gap is largest: per-node breakdown

The 24 evidence-free decision nodes under strict concordance cluster at three structural positions (Figure 3): (i) **NSCLC post-osimertinib** (N7 amivantamab+chemo, N8 HER3-ADC, N9 TROP2-ADC, N10 platinum-doublet salvage, N21 MET-amp salvage, N23 post-amivantamab, N24 post-osi+post-chemo) — 7 of 17 EGFR-mutant nodes are post-osimertinib and 5 are evidence-free at strict; (ii) **NSCLC 1L EGFR-mutant alternative pathways** (N3 amivantamab+lazertinib, N4 second-gen TKI, N5 anti-VEGF+TKI, N11 EGFR ex20ins 1L, N12 EGFR ex20ins post-chemo, N25 high-risk subgroup); (iii) **mBC post-CDK4/6i salvage and special populations** (G9 no-actionable, G13 post-CDK4/6i+post-endo everolimus, G14 post-CDK4/6i+post-chemo SG, G17 visceral-crisis, G22/G24 indolent-disease, G25 post-everolimus, G21 datopotamab HER2-low).

3.5 Operationalization discordance carries across tumors

The mBC prior-CDK4/6i inclusion variable showed substantial cross-trial discordance (ODI 0.64, 95% CI 0.62–0.66 across 3,160 trial pairs in v2). The NSCLC prior-EGFR-TKI inclusion variable shows analogous heterogeneity (some trials require prior 1st/2nd-gen TKI, others require prior osimertinib,




Figure 3. Per-node trial-edge support across the pooled 50-node decision tree, separated into mBC (left) and NSCLC (right). Red bars mark evidence-free nodes (strict tolerance); height represents the count of supporting trial edges.

others permit either) and contributes to the post-osimertinib state-coding ambiguity that drives several evidence-free flags.

4 Discussion

Headline. At adequately-powered multi-tumor scale (50 pooled guideline nodes across mBC HR+/HER2- and NSCLC EGFR + ALK), the pre-registered exact-binomial test rejected $H_0 : \text{EFDPR} \leq 0.25$ at $P = 0.0004$. Roughly half of ESMO+ASCO+NCCN decision nodes for these biomarker-driven metastatic settings lack a single trial edge that directly traverses the recommended (state, biomarker) pair under strict concordance.

NSCLC EGFR post-osimertinib is the largest concentrated gap. Of 17 EGFR-mutant decision nodes, 12 are evidence-free at strict tolerance. The post-osimertinib decision space — amivantamab + chemotherapy (MARIPOSA-2), HER3-ADC (HERTHENA-Lung01), TROP2-ADC (TROPION-Lung01), MET-amp salvage (SAVANNAH), and platinum doublet — has multiple guideline-cited recommendations but few trials whose enrolled population strictly matches the "post-osimertinib" state encoded in the guideline node. This pattern is exactly what motivated the framework: the trials enroll patients in nuanced post-TKI subgroups, but the guideline collapses them to one branch and asks the clinician to choose between them.

NSCLC ALK is the comparator that doesn't show the gap. ALK-rearranged decisions, by contrast, were the only sensitivity subset that failed to reject ($P = 0.63$). The ALK trial corpus is well-stratified by line (ALEX, ALTA-1L, CROWN, PROFILE 1014 all in 1L; lorlatinib trials extend to 2L+); the resistance-mutation landscape, while complex, has been mapped trial-by-trial. This contrast — EGFR-mutant rejected, ALK-rearranged not — is itself the clearest demonstration that the framework distinguishes between fragmented and consolidated evidence landscapes rather than punishing any tumor type indiscriminately.

What this means for guideline committees. The 24 evidence-free decision nodes are not all equally problematic. Some (e.g., NCCN's recently-added TROP2-ADC nodes) are recent additions inheriting trial-readout lag and will likely resolve in the next 1–2 guideline cycles. Others (post-CDK4/6i+post-endo everolimus salvage; post-osimertinib post-chemo) reflect structural gaps in the

trial agenda that the next generation of pivotal trials should target with enrichment populations matching the guideline-node patient state.

Three caveats. First, ESCAT-aligned and liberal tolerance gave $P = 0.10$ on the pooled denominator. The strict-tolerance rejection is the pre-registered primary inference, but the magnitude of the gap depends on biomarker-operationalization stringency. Second, the LLM-extraction pipeline’s pre-adjudication Cohen’s κ was 0.78 for NSCLC and 0.67 for mBC, with documented adjudication trail. Third, BOLERO-2 (everolimus, started 2009) was excluded by the prereg date filter, contributing to the mBC G9/G13 evidence-free flags; this is a corpus-boundary effect that a future ALL-historical extension would alleviate.

Three positives. First, the framework is now empirically validated across two biomarker-driven solid tumors with a pre-registered primary test rejection. Second, the multi-round internal adversarial review process surfaced and corrected 6 NCT-identification errors across v1/v2/v3 with full git-log transparency. Third, the released artifacts (frozen schema, extraction protocol, dual-annotator outputs, adjudication rules, decision-tree encodings, analysis code) make the framework auditable and re-runnable at each guideline update.

Limitations. The 261-trial pooled corpus is a snapshot at freeze 2026-05-16. Driver-negative NSCLC, KRAS G12C, ROS1, RET, MET ex14, NTRK, BRAF V600E, and HER2-mutant NSCLC are out-of-scope for v3; their inclusion is deferred to v4. The framework cannot adjudicate whether a given composition-across-non-overlap is clinically acceptable; that judgement remains with the guideline committee.

Outlook. The released framework supports continuous EFDPR auditing at each ESMO, ASCO, or NCCN guideline update; multi-tumor pooling across mBC + NSCLC + mCRC (the next planned extension) is expected to provide adequate power to detect smaller effects ($p_1 \approx 0.30$).

Data and code availability

ClinicalTrials.gov queries reproduced in `analysis/v3_01_systematic_search_nsclc.py`; guideline decision-tree encodings, structured trial corpora, dual-annotator extractions and adjudication rules released under MIT licence at <https://github.com/htlin/mbc-evidence-dag-paper> at tag v3.0.0; Zenodo DOI minted at release.

Pre-registration

`docs/prereg-v3.md` at commit 4b5bf1a, committed before any NSCLC outcome data.

Conflicts of interest

None declared.

Funding

No specific funding.

References

- [1] Gennari A, André F, Barrios CH, Cortes J, de Azambuja E, DeMichele A, et al. ESMO Clinical Practice Guideline for the diagnosis, staging and treatment of patients with metastatic breast

cancer (Living Guideline 2024 update). *Annals of Oncology*. 2024;35(2):154-72. ESMO Living Guideline; cited herein as the ESMO 2024 update.

- [2] Gradishar WJ, Moran MS, Abraham J, et al. Breast Cancer, Version 3.2024, NCCN Clinical Practice Guidelines in Oncology. *Journal of the National Comprehensive Cancer Network*. 2024;22(5):331-57.
- [3] Giordano SH, Franzoi MAB, Temin S, et al. Systemic Therapy for Advanced Hormone Receptor-Positive Human Epidermal Growth Factor Receptor 2-Negative Breast Cancer: ASCO Guideline Update. *Journal of Clinical Oncology*. 2022;40(23):2612-35.
- [4] Hendriks LE, Kerr KM, Menis J, et al. Oncogene-addicted metastatic non-small-cell lung cancer: ESMO Clinical Practice Guideline for diagnosis, treatment and follow-up. *Annals of Oncology*. 2023;34(4):339-57.
- [5] Ettinger DS, Wood DE, Aisner DL, et al. NCCN Clinical Practice Guidelines in Oncology: Non-Small Cell Lung Cancer Version 5.2024. *Journal of the National Comprehensive Cancer Network*. 2024;22(4):249-74.