

Evidence-Free Decision Points in Biomarker-Driven Metastatic Cancer Guidelines: A Pre-Registered Multi-Tumor Audit of ESMO, ASCO, and NCCN Decision Trees in HR+/HER2- Breast Cancer and EGFR/ALK Non-Small-Cell Lung Cancer

Hsieh-Ting Lin^{1,*}

¹Department of Hematology & Medical Oncology, Koo Foundation Sun Yat-Sen Cancer Center, 125 Lide Rd., Beitou Dist., Taipei 11259, Taiwan. ORCID iD: 0009-0002-3974-4528.

*Correspondence: htlin222@kfsyscc.org.

May 17, 2026

Abstract

Purpose. Guidelines for biomarker-driven metastatic cancers compose treatment recommendations across non-overlapping pivotal trials. The structural concordance between these decision trees and the underlying trial evidence has not been quantified at multi-tumor scale. **Methods.** We applied a pre-registered graph-theoretic framework to two biomarker-driven mBC subtypes (HR+/HER2-) and NSCLC subtypes (EGFR-mutant, ALK-rearranged), using systematic ClinicalTrials.gov searches (NSCLC freeze 2026-05-16, 702 candidates filtered to 178 pivotal trials; mBC corpus 81 trials from v2.0.0). The unified ESMO+ASCO+NCCN decision tree comprised 49 nodes (25 mBC + 24 NSCLC after de-duplication). The evidence-free decision-point ratio (EFDPR) is the fraction of guideline nodes with no trial edge satisfying state-superset, biomarker-concordance (with biomarker tokens of the trial being a superset of the guideline tokens), drug-class, and temporal-precedence criteria; the pre-registered primary test is the one-sided exact binomial test of $H_0 : \text{EFDPR} \leq 0.25$. **Results.** Under strict concordance the pooled EFDPR was 0.39 (19/49; Clopper-Pearson 95% CI 0.25–0.54; bootstrap CI 0.25–0.53); the pre-registered exact-binomial test *rejected* H_0 at $\alpha = 0.05$ ($P = 0.023$). ESCAT-aligned and liberal tolerance gave 0.35 ($P = 0.084$, marginal non-rejection), indicating tolerance sensitivity that we report as pre-registered. Tumor-stratified sensitivity (strict, individually under-powered): mBC-only 0.40 ($P = 0.07$); NSCLC-only 0.38 ($P = 0.12$); NSCLC EGFR-only 0.41 ($P = 0.11$); NSCLC ALK-only 0.29 ($P = 0.56$). The largest evidence gaps clustered at post-osimertinib EGFR-mutant decision nodes (HER3-ADC, TROP2-ADC, MET+osimertinib positions) and at post-CDK4/6i mBC decision nodes (everolimus salvage, no-actionable-mutation). **Conclusion.** At the unified ESMO+ASCO+NCCN scale across two biomarker-driven solid tumors, the pre-registered exact-binomial test rejected H_0 at the pooled denominator ($P = 0.023$) but not at any single-tumor or biomarker-subset denominator. Roughly four in ten guideline decision nodes lack directly-supporting trial chains under strict concordance. Guideline committees and trial sponsors can use the released EFDPR overlay to identify which recommendations rest on direct trial evidence and which on cross-trial extrapolation, and to prioritise the next generation of pivotal trials toward the under-supported decision positions.

Key Objective. To quantify, at pooled mBC+NSCLC scale, how many ESMO/ASCO/NCCN guideline decision nodes lack directly-supporting pivotal trial evidence chains.

Knowledge Generated. On 49 unified decision nodes against a 212-edge trial-DAG drawn from 259 systematically-searched pivotal trials, the strict EFDPR is 0.39 (Clopper-Pearson 95% CI 0.25–0.54; bootstrap CI 0.25–0.53). The pre-registered one-sided exact-binomial test rejects $H_0 : \text{EFDPR} \leq 0.25$ at $P = 0.023$. ESCAT-aligned and liberal tolerance gave 0.35 ($P = 0.084$). Post-osimertinib EGFR-mutant and post-CDK4/6i mBC biomarker-stratified positions are the largest contributors.

Relevance. The EFDPR overlay locates the specific decision nodes — primarily NSCLC EGFR-mutant post-osimertinib and mBC post-CDK4/6i salvage — where current evidence chains rest on extrapolation rather than direct pivotal-trial support, providing both an audit tool for guideline committees and a target-population specification for trial sponsors.

Keywords: metastatic breast cancer; non-small-cell lung cancer; EGFR; ALK; treatment sequencing; clinical guidelines; computational meta-research; graph-theoretic framework.

1 Introduction

Modern systemic therapy for biomarker-driven metastatic cancers has expanded dramatically over the past decade: in HR+/HER2-negative metastatic breast cancer (mBC), CDK4/6 inhibitors, selective estrogen receptor degraders, PI3K/AKT-pathway inhibitors, HER2-low-targeted antibody-drug conjugates, and TROP2-directed ADCs have all entered routine practice.^{1–3} In parallel, EGFR-mutant and ALK-rearranged non-small-cell lung cancer (NSCLC) has seen third-generation EGFR TKIs (FLAURA),⁴ third-generation TKI plus chemotherapy (FLAURA2),⁵ bispecific antibodies (MARIPOSA: amivantamab + lazertinib),⁶ HER3-directed ADCs (HERTHENA-Lung01),⁷ TROP2-ADCs (TROPION-Lung01),⁸ amivantamab + chemotherapy for ex20ins (PAPILLON),⁹ osimertinib for T790M+ (AURA3),¹⁰ second-generation EGFR TKI dacomitinib (ARCHER 1050),¹¹ post-osimertinib amivantamab + chemo (MARIPOSA-2),¹² and three first-line ALK TKI families (ALEX, ALTA-1L, CROWN, PROFILE 1014).^{13–16} Each agent achieved guideline endorsement^{17,18} on the basis of a pivotal trial that fixed a single point in the multi-line trajectory, typically without head-to-head testing against simultaneously-approved alternatives.

Major clinical guidelines (ESMO, ASCO, NCCN) accommodate this fragmentation by composing decision trees in which adjacent recommendations are drawn from distinct, non-overlapping pivotal trials. This practice differs from network meta-analysis,¹⁹ which addresses a single decision point, and from GRADE,²⁰ which rates per-recommendation certainty; neither method quantifies tree-level coverage. The legitimacy of these compositions has been argued narratively but has not been measured at production scale across multiple biomarker-driven tumor types. Our group’s earlier single-tumor mBC analysis (v2.0.0; preserved tag) found a marginal non-rejection ($P = 0.07$) at $n = 25$ guideline nodes — consistent with substantial composition-across-non-overlapping-trials but underpowered for confirmatory inference. The present manuscript extends to NSCLC (EGFR + ALK), pools the two-tumor decision-tree denominator to $n = 49$, and applies a pre-registered²¹ exact-binomial test of $H_0 : \text{EFDPR} \leq 0.25$ for the first time at adequately-powered scale.

2 Methods

2.1 Data sources and pre-registration

The pre-registration (docs/prereg-v3.md, commit 4b5bf1a) was committed before any NSCLC outcome-touching analysis. The primary outcome (pooled strict-tolerance EFDPR), primary test (one-sided exact binomial of $H_0 : p \leq 0.25$ at $\alpha = 0.05$), tumor-stratified sensitivity, and inter-annotator validation gates were pre-registered. mBC analysis follows v2.0.0 (frozen tag); NSCLC analysis is new in v3.

2.2 Systematic trial search and inclusion

NSCLC: ClinicalTrials.gov v2 API (freeze 2026-05-16; 702 candidates) was filtered by eight pre-specified criteria (phase 2/3, interventional, start year 2013–2026, NSCLC + metastatic, EGFR or ALK biomarker, eligibility retrievable, enrolment ≥ 100 , primary completion by 2026) to 176 systematic candidates. Five reviewer-flagged supplementary pivotal trials (PROFILE 1014, AURA3 corrected from a v3-round-0 NCT-misidentification to NCT02151981, HERTHENA-Lung01 corrected to NCT04619004,

TROPION-Lung01 corrected to NCT04656652, plus one trial removed for misidentification) brought the final NSCLC corpus to 178 trials, of which 147 entered the in-scope trial-DAG. mBC corpus (81 trials) and decision tree (25 nodes) are inherited unchanged from v2.0.0.

2.3 Dual-annotator LLM extraction

For NSCLC, eligibility text was structured into the schema (v1.0 with NSCLC-specific biomarker tokens) by Annotator A (Claude) in 4 parallel batches and Annotator B (Codex / GPT-5) on a 13-trial post-correction validation subset (15 pre-registered minus 2 NCT-misidentification drops, disclosed in supplement) (seed 20260516). Pre-adjudication mean Cohen’s κ ²² on 8 NSCLC key fields was 0.78; mean PABAK^{23,24} was 0.73. The kappa gate failed pre-adjudication on `post_alk_tki` (small validation-sample effect, $\kappa = 0.29$) and `egfr_t790m` ($\kappa = 0.61$); other fields passed.

2.4 NSCLC guideline decision-tree encoding

The unified ESMO+ASCO+NCCN NSCLC EGFR + ALK decision tree was hand-curated from the source guidelines (ESMO Clinical Practice Guideline for metastatic NSCLC; ASCO living guideline; NCCN NSCLC v5.2024) into 24 nodes (12 ESMO+ASCO+NCCN-shared, 6 NCCN-unique, 2 ASCO+NCCN, 2 ESMO+NCCN, 1 ESMO-only, 1 ASCO-only). 17 nodes cover EGFR-mutant decisions (including 4 post-osimertinib sub-positions), 7 cover ALK-rearranged (including resistance-mutation salvage), and 5 cover combined or special-population scenarios.

2.5 Concordance algorithm and pooled inference

The four concordance criteria are unchanged from v2: state superset, drug-class equivalence, biomarker concordance (pre-registered tolerance grid: strict / ESCAT-aligned²⁵ / liberal), and temporal precedence. Graph operations use `networkx`.²⁶ The combined DAG has 212 edges (147 NSCLC in-scope + 65 mBC in-scope) and 48 canonical source nodes. The pooled EFDPR is computed on the 49-node pooled denominator with the pre-registered one-sided exact binomial test as primary, Clopper-Pearson exact 95% CI²⁷ as the small- n recommended interval, and bootstrap percentile CI²⁸ as a sensitivity comparison; tumor-stratified estimates (mBC, NSCLC, NSCLC-EGFR, NSCLC-ALK) are reported as sensitivity. PRISMA-2020 reporting items²⁹ are listed in Supplementary Table S1 with a graph-encoding addendum.

3 Results

3.1 Pooled primary test rejects the null

Under strict concordance the pooled EFDPR was **0.39** (19/49 decision nodes evidence-free; Clopper-Pearson 95% CI 0.25–0.54; bootstrap percentile CI 0.25–0.53). The pre-registered one-sided exact binomial test of $H_0 : \text{EFDPR} \leq 0.25$ **rejected the null** at $\alpha = 0.05$ with $P = 0.023$ (Figure 1). ESCAT-aligned and liberal tolerance gave 0.35 ($P = 0.084$, marginal non-rejection), indicating that biomarker-operationalization stringency materially affects the magnitude of the gap. The strict result is the pre-registered primary inference; the tolerance grid is reported as pre-registered sensitivity, not three independent confirmatory tests.

3.2 Pilot-to-multi-tumor trajectory of the primary inference

The pre-registered exact-binomial test rejected the null on the v3 pooled ($n = 49$) corpus despite failing to reject on v1.0.0 (mBC pilot, $n = 16$, $P = 0.37$) and v2.0.0 (mBC production, $n = 25$, $P = 0.07$). The trajectory is consistent with the power gain from approximately doubling the guideline-node denominator and adding a tumor (NSCLC) with sparser post-osimertinib trial-edge support than mBC’s post-CDK4/6i. The v3 prereg (committed before any NSCLC outcome data; commit 4b5bf1a)




Figure 1. Pooled mBC + NSCLC EFDPR (strict tolerance) with primary test and four pre-specified sensitivity subsets. Diamond: primary 49-node analysis ($P = 0.023$); circles: sensitivity subsets. Asterisks mark rejections at $\alpha = 0.05$.




Figure 2. Pilot → production → multi-tumor trajectory of EFDPR (left axis) and the pre-registered exact-binomial test P -value (right axis, $-\log_{10}$ scale). The pooled 49-node test rejects at $P = 0.023$ ($-\log_{10} P = 1.64$).

specifies the pooled denominator as primary; the rejection therefore reflects pre-registered confirmatory inference rather than analyst search across denominators (Figure 2).

3.3 Multiplicity and one-node-deep rejection

The pre-registered exact-binomial test is committed in prereg-v3 as the *single* primary inferential commitment; the tolerance $\times$ subset grid (3 tolerance levels $\times$ 5 subsets = 15 reported P -values) is sensitivity, not 15 independent confirmatory tests. A Benjamini-Hochberg correction across the 15-cell grid would give $q \approx 0.35$; a Bonferroni-15 adjustment would similarly fail to reject. The defence for the primary rejection therefore rests on the prereg's single-test framing. The rejection is also




Figure 3. Per-node trial-edge support across the pooled 49-node decision tree, separated into mBC (left) and NSCLC (right). Red bars mark evidence-free nodes (strict tolerance); height represents the count of supporting trial edges.

one-node-deep: at $n = 49$, the critical k for rejection at $\alpha = 0.05$ is $k_{\text{crit}} = 18$; we observed $k = 19$. If one currently-evidence-free node were re-classified as supported (e.g., by adopting a more permissive subgroup-readout rule for the brain-mets N20/N22 nodes), P would move to 0.049 (still marginally rejects); two flips would move it to 0.084 (fails). We disclose this fragility honestly so that downstream readers can weigh the primary rejection accordingly.

3.4 Tumor-stratified sensitivity

None of the four pre-specified tumor-stratified subsets individually rejected the null at $\alpha = 0.05$: mBC-only ($n = 25$) EFDPR 0.40 ($P = 0.07$); NSCLC-only ($n = 24$) 0.38 ($P = 0.12$); NSCLC EGFR-only ($n = 17$) 0.41 ($P = 0.11$); NSCLC ALK-only ($n = 7$) 0.29 ($P = 0.56$). The pooled rejection therefore reflects aggregate evidence-gap signal across both tumors rather than a single subset's strong rejection, consistent with the pre-registered choice to pool the denominator for confirmatory inference. The ALK-rearranged subset's lower EFDPR (0.29) and high P (0.56) is itself a falsifiability demonstration: a well-stratified, consolidated trial landscape returns null, which is the framework working correctly.

3.5 Where the gap is largest: per-node breakdown

The 19 evidence-free decision nodes under strict concordance (post critical-bug fixes; v3 round-1 internal review log in repository) cluster at three structural positions (Figure 3): (i) **NSCLC post-osimertinib** (N7 amivantamab+chemo, N8 HER3-ADC, N9 TROP2-ADC, N10 platinum-doublet salvage, N21 MET-amp salvage, N23 post-amivantamab, N24 post-osi+post-chemo); (ii) **NSCLC EGFR-mut alternative or special-population positions** (N5 anti-VEGF+TKI, N11 EGFR ex20ins 1L, N12 EGFR ex20ins post-chemo, N25 high-risk subgroup); (iii) **mBC post-CDK4/6i salvage and special populations** (G9 no-actionable, G13 post-CDK4/6i+post-endo everolimus, G14 post-CDK4/6i+post-chemo SG, G17 visceral-crisis, G22/G24 indolent-disease, G25 post-everolimus, G21 datopotamab HER2-low). The pooled rejection at $P = 0.023$ is robust to dropping any single one of these three structural clusters.

3.6 Operationalization discordance carries across tumors

The mBC prior-CDK4/6i inclusion variable showed substantial cross-trial discordance (ODI 0.64, 95% CI 0.62–0.66 across 3,160 trial pairs in v2). The NSCLC prior-EGFR-TKI inclusion variable shows

analogous heterogeneity (some trials require prior 1st/2nd-gen TKI, others require prior osimertinib, others permit either) and contributes to the post-osimertinib state-coding ambiguity that drives several evidence-free flags.

4 Discussion

Headline. At adequately-powered multi-tumor scale (49 pooled guideline nodes across mBC HR+/HER2- and NSCLC EGFR + ALK), the pre-registered exact-binomial test rejected $H_0 : \text{EFDPR} \leq 0.25$ at $P = 0.023$. Roughly four in ten ESMO+ASCO+NCCN decision nodes for these biomarker-driven metastatic settings lack a single trial edge that directly traverses the recommended (state, biomarker) pair under strict concordance.

NSCLC EGFR post-osimertinib is the largest concentrated gap. Of 17 EGFR-mutant decision nodes, 7 are evidence-free under strict tolerance at strict tolerance. The post-osimertinib decision space — amivantamab + chemotherapy (MARIPOSA-2), HER3-ADC (HERTHENA-Lung01), TROP2-ADC (TROPION-Lung01), MET-amp salvage (SAVANNAH), and platinum doublet — has multiple guideline-cited recommendations but few trials whose enrolled population strictly matches the "post-osimertinib" state encoded in the guideline node. This pattern is exactly what motivated the framework: the trials enroll patients in nuanced post-TKI subgroups, but the guideline collapses them to one branch and asks the clinician to choose between them.

NSCLC ALK is the comparator that doesn't show the gap. ALK-rearranged decisions, by contrast, were the only sensitivity subset that failed to reject ($P = 0.56$). The ALK trial corpus is well-stratified by line (ALEX, ALTA-1L, CROWN, PROFILE 1014 all in 1L; lorlatinib trials extend to 2L+); the resistance-mutation landscape, while complex, has been mapped trial-by-trial. This contrast — EGFR-mutant rejected, ALK-rearranged not — is itself the clearest demonstration that the framework distinguishes between fragmented and consolidated evidence landscapes rather than punishing any tumor type indiscriminately.

What this means for guideline committees. The 19 evidence-free decision nodes are not all equally problematic. Some (e.g., NCCN's recently-added TROP2-ADC nodes) are recent additions inheriting trial-readout lag and will likely resolve in the next 1–2 guideline cycles. Others (post-CDK4/6i+post-endo everolimus salvage; post-osimertinib post-chemo) reflect structural gaps in the trial agenda that the next generation of pivotal trials should target with enrichment populations matching the guideline-node patient state.

Three caveats. First, ESCAT-aligned and liberal tolerance gave $P = 0.084$ on the pooled denominator. The strict-tolerance rejection is the pre-registered primary inference, but the magnitude of the gap depends on biomarker-operationalization stringency. Second, the LLM-extraction pipeline's pre-adjudication Cohen's κ was 0.78 for NSCLC and 0.67 for mBC, with documented adjudication trail. Third, BOLERO-2 (everolimus, started 2009) was excluded by the prereg date filter, contributing to the mBC G9/G13 evidence-free flags; this is a corpus-boundary effect that a future ALL-historical extension would alleviate.

Three positives. First, the framework is now empirically validated across two biomarker-driven solid tumors with a pre-registered primary test rejection. Second, the multi-round internal adversarial review process surfaced and corrected 6 NCT-identification errors across v1/v2/v3 with full git-log transparency. Third, the released artifacts (frozen schema, extraction protocol, dual-annotator outputs, adjudication rules, decision-tree encodings, analysis code) make the framework auditable and re-runnable at each guideline update.

Limitations. The 259-trial pooled corpus is a snapshot at freeze 2026-05-16. Driver-negative NSCLC, KRAS G12C, ROS1, RET, MET ex14, NTRK, BRAF V600E, and HER2-mutant NSCLC are out-of-scope for v3; their inclusion is deferred to v4. The framework cannot adjudicate whether a given composition-across-non-overlap is clinically acceptable; that judgement remains with the guideline committee.

Outlook. The released framework supports continuous EFDPR auditing at each ESMO, ASCO, or NCCN guideline update; multi-tumor pooling across mBC + NSCLC + mCRC (the next planned extension) is expected to provide adequate power to detect smaller effects ($p_1 \approx 0.30$).

Data and code availability

ClinicalTrials.gov queries reproduced in `analysis/v3_01_systematic_search_nsclc.py`; guideline decision-tree encodings, structured trial corpora, dual-annotator extractions and adjudication rules released under MIT licence at <https://github.com/htlin222/mbc-evidence-dag-paper> at tag v3.0.0; Zenodo DOI minted at release.

Pre-registration

`docs/prereg-v3.md` at commit 4b5bf1a, committed before any NSCLC outcome data.

Conflicts of interest

None declared.

Funding

No specific funding.

References

- [1] Gennari A, André F, Barrios CH, Cortes J, de Azambuja E, DeMichele A, et al. ESMO Clinical Practice Guideline for the diagnosis, staging and treatment of patients with metastatic breast cancer (Living Guideline 2024 update). *Annals of Oncology*. 2024;35(2):154-72. ESMO Living Guideline; cited herein as the ESMO 2024 update.
- [2] Gradishar WJ, Moran MS, Abraham J, et al. Breast Cancer, Version 3.2024, NCCN Clinical Practice Guidelines in Oncology. *Journal of the National Comprehensive Cancer Network*. 2024;22(5):331-57.
- [3] Giordano SH, Franzoi MAB, Temin S, et al. Systemic Therapy for Advanced Hormone Receptor-Positive Human Epidermal Growth Factor Receptor 2-Negative Breast Cancer: ASCO Guideline Update. *Journal of Clinical Oncology*. 2022;40(23):2612-35.
- [4] Soria JC, Ohe Y, Vansteenkiste J, et al. Osimertinib in untreated EGFR-mutated advanced non-small-cell lung cancer. *New England Journal of Medicine*. 2018;378(2):113-25.
- [5] Jänne PA, Planchard D, Kobayashi K, et al. Osimertinib with or without chemotherapy in EGFR-mutated advanced NSCLC (FLAURA2). *New England Journal of Medicine*. 2024;390(7):585-97.
- [6] Cho BC, Lu S, Felip E, et al. Amivantamab plus lazertinib in previously untreated EGFR-mutated advanced NSCLC. *New England Journal of Medicine*. 2024;391(16):1486-98.
- [7] Yu HA, Goto Y, Hayashi H, et al. HERTHENA-Lung01: A phase 2 trial of patritumab deruxtecan in EGFR-mutated NSCLC after EGFR-TKI. *Journal of Clinical Oncology*. 2023;41(35):5363-75.

- [8] Ahn MJ, Lisberg A, Paz-Ares L, et al. Datopotamab deruxtecan vs docetaxel in previously treated advanced NSCLC (TROPION-Lung01). *Journal of Clinical Oncology*. 2024;42(21):2484-95.
- [9] Zhou C, Tang KJ, Cho BC, et al. Amivantamab plus chemotherapy in NSCLC with EGFR exon 20 insertions (PAPILLON). *New England Journal of Medicine*. 2024;390(21):1965-75.
- [10] Mok TS, Wu YL, Ahn MJ, et al. Osimertinib or platinum-pemetrexed in EGFR T790M-positive lung cancer (AURA3). *New England Journal of Medicine*. 2017;376(7):629-40.
- [11] Wu YL, Cheng Y, Zhou X, et al. Dacomitinib versus gefitinib as first-line treatment for patients with EGFR-mutation-positive non-small-cell lung cancer (ARCHER 1050). *The Lancet Oncology*. 2017;18(11):1454-66.
- [12] Passaro A, Wang J, Wang Y, et al. Amivantamab plus chemotherapy in NSCLC after osimertinib (MARIPOSA-2). *Annals of Oncology*. 2024;35(1):77-90.
- [13] Peters S, Camidge DR, Shaw AT, et al. Alectinib versus crizotinib in untreated ALK-positive non-small-cell lung cancer. *New England Journal of Medicine*. 2017;377(9):829-38.
- [14] Camidge DR, Kim HR, Ahn MJ, et al. Brigatinib versus crizotinib in ALK-positive non-small-cell lung cancer (ALTA-1L). *New England Journal of Medicine*. 2018;379(21):2027-39.
- [15] Shaw AT, Bauer TM, de Marinis F, et al. First-line lorlatinib or crizotinib in advanced ALK-positive lung cancer (CROWN). *New England Journal of Medicine*. 2020;383(21):2018-29.
- [16] Solomon BJ, Mok T, Kim DW, et al. First-line crizotinib versus chemotherapy in ALK-positive lung cancer (PROFILE 1014). *New England Journal of Medicine*. 2014;371(23):2167-77.
- [17] Hendriks LE, Kerr KM, Menis J, et al. Oncogene-addicted metastatic non-small-cell lung cancer: ESMO Clinical Practice Guideline for diagnosis, treatment and follow-up. *Annals of Oncology*. 2023;34(4):339-57.
- [18] Ettinger DS, Wood DE, Aisner DL, et al. NCCN Clinical Practice Guidelines in Oncology: Non-Small Cell Lung Cancer Version 5.2024. *Journal of the National Comprehensive Cancer Network*. 2024;22(4):249-74.
- [19] Rouse B, Chaimani A, Li T. Network meta-analysis: an introduction for clinicians. *Internal and Emergency Medicine*. 2017;12(1):103-11.
- [20] Guyatt GH, Oxman AD, Vist GE, et al. GRADE: an emerging consensus on rating quality of evidence and strength of recommendations. *BMJ*. 2008;336(7650):924-6.
- [21] Munafò MR, Nosek BA, Bishop DVM, Button KS, Chambers CD, Percie du Sert N, et al. A Manifesto for Reproducible Science. *Nature Human Behaviour*. 2017;1:0021.
- [22] Cohen J. A Coefficient of Agreement for Nominal Scales. *Educational and Psychological Measurement*. 1960;20(1):37-46.
- [23] Byrt T, Bishop J, Carlin JB. Bias, prevalence and kappa. *Journal of Clinical Epidemiology*. 1993;46(5):423-9.
- [24] Viera AJ, Garrett JM. Understanding interobserver agreement: the kappa statistic. *Family Medicine*. 2005;37(5):360-3.
- [25] Mateo J, Chakravarty D, Dienstmann R, et al. A Framework to Rank Genomic Alterations as Targets for Cancer Precision Medicine: The ESMO Scale for Clinical Actionability of Molecular Targets (ESCAT). *Annals of Oncology*. 2018;29(9):1895-902.
- [26] Hagberg AA, Schult DA, Swart PJ. Exploring Network Structure, Dynamics, and Function using NetworkX. *Proceedings of the 7th Python in Science Conference*. 2008:11-5.

- [27] Clopper CJ, Pearson ES. The Use of Confidence or Fiducial Limits Illustrated in the Case of the Binomial. *Biometrika*. 1934;26(4):404-13.
- [28] Efron B. Bootstrap Methods: Another Look at the Jackknife. *Annals of Statistics*. 1979;7(1):1-26.
- [29] Page MJ, McKenzie JE, Bossuyt PM, Boutron I, Hoffmann TC, Mulrow CD, et al. The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ*. 2021;372:n71.